

# Assignment 01: Evaluate the Ad Budget Dataset of XYZ Firm

The comments/sections provided are your cues to perform the assignment. You don't need to limit yourself to the number of rows/cells provided. You can add additional rows in each section to add more lines of code.

If at any point in time you need help on solving this assignment, view our demo video to understand the different steps of the code.

Happy coding!

## Analysing Ad Budgets for different media channels

### DESCRIPTION

#### Problem:

The given dataset contains ad budgets for different media channels and the corresponding ad sales of XYZ firm. Evaluate the dataset to:

- Find the features or media channels used by the firm
- Find the sales figures for each channel
- Create a model to predict the sales outcome
- Split as training and testing datasets for the model
- Calculate the Mean Square Error (MSE)

#### 1: Import the dataset

```
In [1]: #Import the required libraries
import pandas as pd
```

```
In [2]: #Import the advertising dataset
df_ad_data = pd.read_csv('D:\\NIPUN_SC_REC\\3_Practice_Project\\Course_5_Data Science with Python\\Practice_projects\\advertising.csv')
```

#### 2: Analyze the dataset

```
In [3]: #View the initial few records of the dataset
df_ad_data.head()
```

```
Out[3]:
```

	TV Ad Budget (\$)	Radio Ad Budget (\$)	Newspaper Ad Budget (\$)	Sales (\$)
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

```
In [4]: #Check the total number of elements in the dataset
df_ad_data.size
```

```
Out[4]: 800
```

#### 3: Find the features or media channels used by the firm

```
In [5]: #Check the number of observations (rows) and attributes (columns) in the dataset
df_ad_data.shape
```

```
Out[5]: (200, 4)
```

```
In [6]: #View the names of each of the attributes
df_ad_data.columns
```

```
Out[6]: Index(['TV Ad Budget ($)', 'Radio Ad Budget ($)', 'Newspaper Ad Budget ($)',
              'Sales ($)'],
              dtype='object')
```

#### 4: Create objects to train and test the model; find the sales figures for each channel

```
In [7]: #Create a feature object from the columns
X_features = df_ad_data[['TV Ad Budget ($)', 'Radio Ad Budget ($)', 'Newspaper Ad Budget ($)']]
```

```
In [8]: #View the feature object
X_features.head()
```

```
Out[8]:
```

	TV Ad Budget (\$)	Radio Ad Budget (\$)	Newspaper Ad Budget (\$)
1	230.1	37.8	69.2
2	44.5	39.3	45.1
3	17.2	45.9	69.3
4	151.5	41.3	58.5
5	180.8	10.8	58.4

```
In [9]: #Create a target object (Hint: use the sales column as it is the response of the dataset)
Y_target = df_ad_data[['Sales ($)']]
```

```
In [10]: #View the target object
Y_target.head()
```

```
Out[10]:
```

	Sales (\$)
1	22.1
2	10.4
3	9.3
4	18.5
5	12.9

```
In [11]: #Verify if all the observations have been captured in the feature object
X_features.shape
```

```
Out[11]: (200, 3)
```

```
In [12]: #Verify if all the observations have been captured in the target object
Y_target.shape
```

```
Out[12]: (200, 1)
```

#### 5: Split the original dataset into training and testing datasets for the model

```
In [13]: #Split the dataset (by default, 75% is the training data and 25% is the testing data)
from sklearn.model_selection import train_test_split
#x_train, x_test, y_train, y_test = model_selection.train_test_split(X_features,Y_target,random_state=1)
from sklearn import model_selection
x_train, x_test, y_train, y_test = model_selection.train_test_split(X_features,Y_target,random_state=1)
```

```
In [14]: #Verify if the training and testing datasets are split correctly (Hint: use the shape() method)
print (x_train.shape,y_train.shape, x_test.shape, y_test.shape)
```

```
(150, 3) (150, 1) (50, 3) (50, 1)
```

#### 6: Create a model to predict the sales outcome

```
In [15]: #Create a linear regression model
from sklearn.linear_model import LinearRegression
linReg = LinearRegression()
linReg.fit(x_train,y_train)
```

```
Out[15]: LinearRegression()
```

```
In [16]: #Print the intercept and coefficients
print (linReg.intercept_)
print (linReg.coef_)
```

```
[2.87696662]
[[0.04656457 0.17915812 0.00345046]]
```

```
In [17]: #Predict the outcome for the testing dataset
y_predict=linReg.predict(x_test)
y_predict
```

```
Out[17]: array([[21.70910292],
                [16.41055243],
                [ 7.60955058],
                [17.80769552],
                [18.6146359 ],
                [23.83573998],
                [16.32488681],
                [13.43225536],
                [ 9.17173403],
                [17.333853 ],
                [14.44479482],
                [ 9.83511973],
                [17.18797614],
                [16.73086831],
                [15.05529391],
                [15.61434433],
                [12.42541574],
                [17.17716376],
                [11.08827566],
                [18.00537501],
                [ 9.28438889],
                [12.98458458],
                [ 8.79950614],
                [10.42382499],
                [11.3846456 ],
                [14.98082512],
                [ 9.78853268],
                [19.39643187],
                [18.18099936],
                [17.12807566],
                [21.54670213],
                [14.69809481],
                [16.24641438],
                [12.32114579],
                [19.92422501],
                [15.32498602],
                [13.88726522],
                [10.03162255],
                [20.93105915],
                [ 7.44936831],
                [ 3.64695761],
                [ 7.22020178],
                [ 5.9962782 ],
                [18.43381853],
                [ 8.39408045],
                [14.08371047],
                [15.02195699],
                [20.35836418],
                [20.57036347],
                [19.60636679]])
```

#### 7: Calculate the Mean Square Error (MSE)

```
In [18]: #Import required libraries for calculating MSE (mean square error)
from sklearn import metrics
import numpy as np
```

```
In [19]: #Calculate the MSE(Mean Squared Error)
print(metrics.mean_squared_error(y_test, y_predict))
```

```
1.9730456202283364
```

```
In [20]: #calculating RMSE(Root Mean Squared Error)
print(np.sqrt(metrics.mean_squared_error(y_test, y_predict)))
```

```
1.4046514230328948
```

```
In [21]: #variance Score
print (linReg.score(x_test,y_test))
```

```
0.9156213613792233
```